


Part 5


- Initial Network

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 150, 150, 3)	0
conv2d (Conv2D)	(None, 150, 150, 8)	224
max_pooling2d (MaxPooling2D)	(None, 75, 75, 8)	0
conv2d_1 (Conv2D)	(None, 75, 75, 16)	1,168
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 16)	0
conv2d_2 (Conv2D)	(None, 37, 37, 32)	4,640
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 32)	0
conv2d_3 (Conv2D)	(None, 18, 18, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 64)	0
flatten (Flatten)	(None, 5184)	0

```
| dense (Dense) | (None, 16) | 82,960  
|-----+-----+-----  
| dense_1 (Dense) | (None, 3) | 51  
|-----+-----+-----  
+  
Total params: 107,539 (420.07 KB)  
Trainable params: 107,539 (420.07 KB)  
Non-trainable params: 0 (0.00 B)
```





-
- * Evaluating basic_model
- 30/30 ————— 3s 90ms/step - accuracy: 0.7296 - loss: 0.7036

Part 6

- Initial Network

```
Model: "sequential"
```


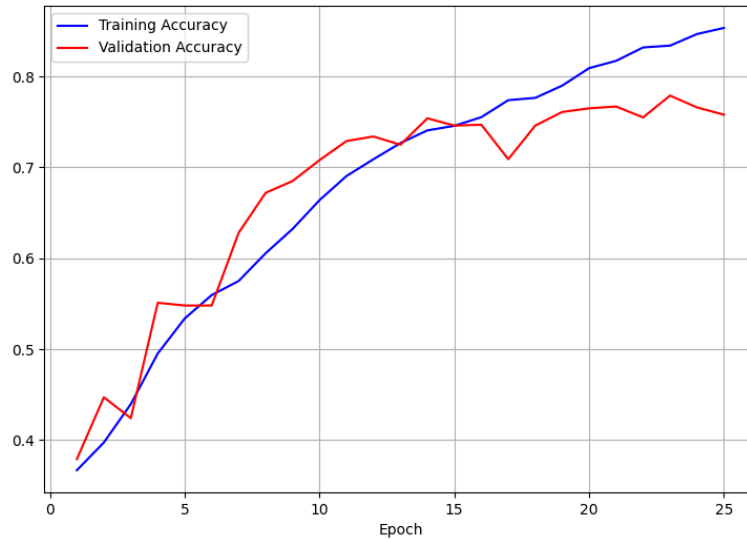
Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 150, 150, 3)	0
conv2d (Conv2D)	(None, 150, 150, 8)	224
max_pooling2d (MaxPooling2D)	(None, 75, 75, 8)	0

conv2d_1 (Conv2D)	(None, 75, 75, 16)	1,168
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 16)	0
conv2d_2 (Conv2D)	(None, 37, 37, 32)	4,640
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 32)	0
conv2d_3 (Conv2D)	(None, 18, 18, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 64)	0
conv2d_4 (Conv2D)	(None, 9, 9, 128)	73,856
max_pooling2d_4 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 16)	32,784
dropout (Dropout)	(None, 16)	0
dense_1 (Dense)	(None, 3)	51

```

+-----+
+
Total params: 131,219 (512.57 KB)
Trainable params: 131,219 (512.57 KB)
Non-trainable params: 0 (0.00 B)

```



```

* Evaluating basic_model
30/30 ----- 3s 85ms/step - accuracy: 0.7398 - loss: 0.6369

```

- I changed the amount of convolutional by adding a final convolutional layer with 128 filters/kernels to the already 4 convolutional layers (8, 16, 32, 64 filters/kernels respectively).
- I also cut the learning rate in half to avoid overfitting.
- I added a dropout layer with 0.3 dropout rate after the one and only dense layer (dense layer has 16 filters/kernels)

Part 7

Game trace:

||||

||||

||||

Turn: X

Player X took position (2, 1).

||||

||||

||X||

Turn: O

reference:

row 0 is neutral.

row 1 is happy.

row 2 is surprise.

1/1 ————— 0s 83ms/step

Emotion detected as happy (row 1). Enter 'text' to use text input instead (0, 1 or 2). Otherwise, press Enter to continue.

reference:

col 0 is neutral.

col 1 is happy.

col 2 is surprise.

1/1 ————— 0s 64ms/step

Emotion detected as neutral (col 0). Enter 'text' to use text input instead (0, 1 or 2). Otherwise, press Enter to continue.

Player O took position (1, 0).

||||

|O||

||X||

Turn: X

Player X took position (0, 0).

|X||

|O||

||X||

Turn: O

reference:

row 0 is neutral.

row 1 is happy.

row 2 is surprise.

1/1 ————— 0s 59ms/step

Emotion detected as happy (row 1). Enter 'text' to use text input instead (0, 1 or 2). Otherwise, press Enter to continue.

reference:

col 0 is neutral.

col 1 is happy.

col 2 is surprise.

1/1 _____ 0s 64ms/step

Emotion detected as happy (col 1). Enter 'text' to use text input instead (0, 1 or 2). Otherwise, press Enter to continue.

Player O took position (1, 1).

|X| | |

|O|O| |

| |X| |

Turn: X

Player X took position (2, 2).

|X| | |

|O|O| |

| |X|X|

Turn: O

reference:

col 0 is neutral.

col 1 is happy.

col 2 is surprise.

1/1 _____ 0s 61ms/step

Emotion detected as neutral (col 0). Enter 'text' to use text input instead (0, 1 or 2). Otherwise, press Enter to continue.

reference:

col 0 is neutral.

col 1 is happy.

col 2 is surprise.

1/1 _____ 0s 62ms/step

Emotion detected as happy (col 1). Enter 'text' to use text input instead (0, 1 or 2). Otherwise, press Enter to continue.

text

2

Player O took position (1, 2).

|X| | |

|O|O|O|

| |X|X|

Player O has won!

Questions

1. How well did your interface work?

It didn't work as anticipated, but worked enough to win the game.

2. Did it recognize your facial expressions with the same accuracy as it achieved against the test set?

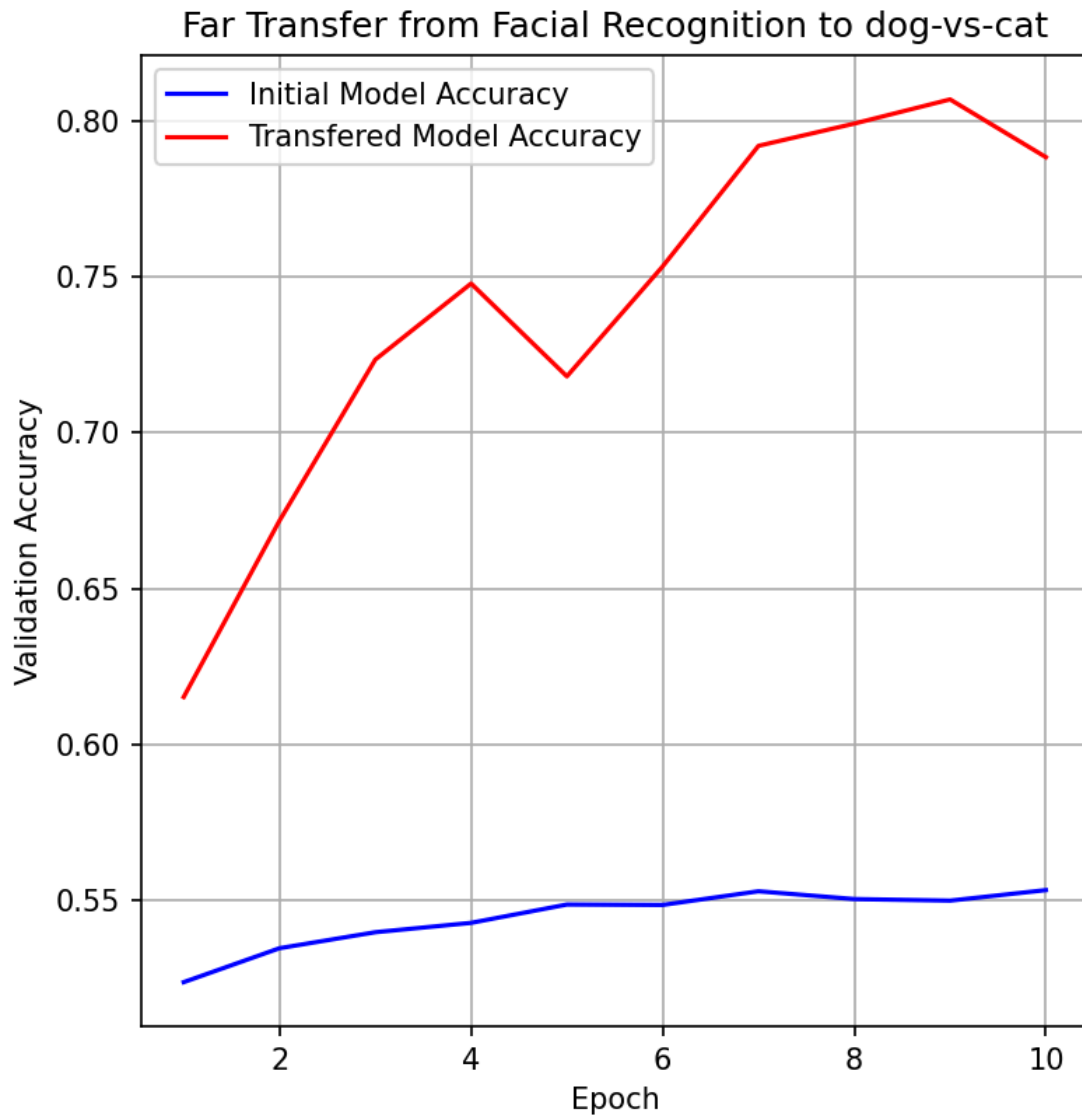
Not really.

3. If not, why not?

We spent some time trying to understand why the bot could somewhat recognize neutral and happy expressions but struggled completely with surprised faces. After repeatedly swapping out the base model it loaded, we deduced that the bot likely "cheats" to achieve the desired results in other steps. To reach 80% accuracy, it doesn't actually matter much what it selects for surprised faces, assuming the test dataset is evenly distributed. This is because if it correctly identifies neutral and happy faces nearly 100% of the time, it wouldn't need to adjust the base probability of random selection significantly to maintain 80% accuracy overall. The math behind this is: $\{\frac{1}{3} + \frac{1}{3} + (\frac{1}{3} * \frac{1}{3}) = \frac{7}{9} = 0.7777...\}$ This means the bot only needs to be about 3% more accurate than random guessing to achieve 80% accuracy.

Code for `_get_emotion`

```
def _get_emotion(self, img) -> int:
    model = models.load_model('TicTacToeModel/game_model.keras')
    img_resized = cv2.resize(img, tuple(image_size))
    img_resized = img_resized / 255.0
    img_resized = np.repeat(img_resized[:, :, np.newaxis], 3, axis=-1)
    img_resized = np.expand_dims(img_resized, axis=0)
    predictions = model.predict(img_resized)
    emotion = int(np.argmax(predictions))
    return emotion
```

Final Accuracy

Restoring model weights from the end of the best epoch: 10.

* Evaluating transfered_model

98/98 ————— 5s 45ms/step - accuracy: 0.5108 - loss: 1.3916

Restoring model weights from the end of the best epoch: 9.

* Evaluating random_model

98/98 ————— 6s 61ms/step - accuracy: 0.4762 - loss: 1.9867

Facial recognition image categories



happy



neutral



surprise

Dog-Vs-Cat image categories



dog



cat

Summaries

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 150, 150, 3)	0
conv2d (Conv2D)	(None, 150, 150, 8)	224
max_pooling2d (MaxPooling2D)	(None, 75, 75, 8)	0
conv2d_1 (Conv2D)	(None, 75, 75, 16)	1,168
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 16)	0
conv2d_2 (Conv2D)	(None, 37, 37, 32)	4,640
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 32)	0
conv2d_3 (Conv2D)	(None, 18, 18, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 64)	0
conv2d_4 (Conv2D)	(None, 9, 9, 128)	73,856
max_pooling2d_4 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 16)	32,784
dropout (Dropout)	(None, 16)	0
dense_1 (Dense)	?	0 (unbuilt)
dropout_1 (Dropout)	?	0
output_layer (Dense)	?	0 (unbuilt)

Total params: 131,168 (512.38 KB)

Trainable params: 0 (0.00 B)

Non-trainable params: 131,168 (512.38 KB)

Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 150, 150, 3)	0
conv2d (Conv2D)	(None, 150, 150, 8)	224
max_pooling2d (MaxPooling2D)	(None, 75, 75, 8)	0
conv2d_1 (Conv2D)	(None, 75, 75, 16)	1,168
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 16)	0
conv2d_2 (Conv2D)	(None, 37, 37, 32)	4,640
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 32)	0
conv2d_3 (Conv2D)	(None, 18, 18, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 64)	0
conv2d_4 (Conv2D)	(None, 9, 9, 128)	73,856
max_pooling2d_4 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 16)	32,784
dropout (Dropout)	(None, 16)	0
dense_2 (Dense)	?	0 (unbuilt)
dropout_2 (Dropout)	?	0
output_layer_2 (Dense)	?	0 (unbuilt)

Total params: 131,168 (512.38 KB)

Trainable params: 131,168 (512.38 KB)

Non-trainable params: 0 (0.00 B)